

PROJECT 4.1
BREAST CANCER – CLASSIFICATION MODEL

SCRIPT

```
setwd("C:/Users/tsraj/Desktop/Acadgild students projects/project4")

library(readr)

CancerData <- read_csv("CancerData.csv")

print(paste("rows:", nrow(df), "cols:", ncol(CancerData)))

View(CancerData)

summary(CancerData)

dim(CancerData)

names(CancerData)

#CancerData<- CancerData[-1]

CancerData$diagnosis <- factor(CancerData$diagnosis, levels = c("B", "M"),
                              labels = c("Benign", "Malignant"))

names(CancerData)

library(mice)

library(readr,dplyr)

library("ggplot2")

library("corrplot")

library("gridExtra")

library("pROC")

library("MASS")

library("caTools")

library("caret")

library(randomForest)

library(rpart)

library(rpart.plot)

library(rattle)

library(ggplot2)

library(Amelia)
```

```

library(class)

library(gmodels)

misssmap(CancerData, main="Missing Data Map", col=c("#FF4081", "#3F51B5"),
         legend=FALSE)

data<-CancerData

data[,33]<-NULL

barplot(table(data$diagnosis), xlab = "Type of tumor", ylab="Numbers per type")

str(data)

any(is.na(data))

# visualize the missing values using the missing map from the Amelia package
misssmap(data,col=c("yellow","red"))

data$diagnosis<-as.factor(data$diagnosis)

summary(data)

qplot(radius_mean, data=data, colour=diagnosis, geom="density",
      main="Radius mean for each tumor type")

qplot(smoothness_mean, data=data, colour=diagnosis, geom="density",
      main="Smoothness mean for each tumor type")

qplot(concavity_mean, data=data, colour=diagnosis, geom="density",
      main="Concavity mean for each tumor type")

qplot(area_worst , data=data, colour=diagnosis, geom="density",
      main="area worst for each tumor type")

# Looking at distribution for area.mean variable

plot.new()

hist(CancerData$area_mean,
     main = 'Distribution of Cell Area Means',
     xlab = 'Mean Area',
     col = 'green')

#we find that the data is imbalanced and also there is a lot of corelation between the attributes

```

```

## we find that there are no missing values

## we find that data is little unbalanced

prop.table(table(data$diagnosis))

## we then show some correlation

corr_mat<-cor(data[,3:ncol(data)])

corrplot(corr_mat)


plot.new()

plot(data$area_mean ~data$concavity_mean)

title('Basic Scatterplot')

ggplot(data, aes(x=data$area_worst)) + geom_histogram(binwidth = 1, fill = "yellow", color =
"black")

ggplot(data, aes(x=data$area_mean)) + geom_histogram(binwidth = 1, fill = "green", color = "red")


#Modelling

#We are going to get a training and a testing set to use when building some models:

set.seed(1234)

data_index<-createDataPartition(data$diagnosis,p=0.75,list = FALSE)

train_data<-data[data_index,-1]

test_data<-data[-data_index,-1]


## Applying learning models

fitControl <- trainControl(method="cv",
                           number = 5,
                           preProcOptions = list(thresh = 0.99), # threshold for pca preprocess
                           classProbs = TRUE,
                           summaryFunction = twoClassSummary)

#Model1: Random Forest

#Building the model on the training data

## random forest

model_rf <- train(diagnosis~.,

```

```

train_data,
method="ranger",
metric="ROC",
#tuneLength=10,
#tuneGrid = expand.grid(mtry = c(2, 3, 6)),
preProcess = c('center', 'scale'),
trControl=fitControl)

#Testing on the testing data
## testing for random forests
pred_rf <- predict(model_rf, test_data)
cm_rf <- confusionMatrix(pred_rf, test_data$diagnosis, positive = "M")
cm_rf
# We find the accuracy of the model is 100%
#Random forest model- takes decision trees and averages them
normalize<-function(x){return((x-min(x))/(max(x)-min(x)))}
data$diagnosis<-as.numeric(data$diagnosis)
data_n<-as.data.frame(lapply(data,normalize))
traindata_n<--data_n[1:426,]
testdata_n<-data_n[427:569,]
rf <- randomForest(diagnosis ~., data= traindata_n, ntree =300, mtry = 5, importance = TRUE)
print(rf)
plot.new()
varImpPlot(rf, type = 1, pch =8, col = 2, cex =0.8, main = "cancerdata")
abline(v= 45, col= "red")
library(party)
#cf1 <- cforest(diagnosis ~ . , data=traindata_n , control=fitControl(mtry=5,ntree=300)) # fit the
random forest

#varimp(cf1) # get variable importance, based on mean decrease in accuracy

```

```
#varimp(cf1, conditional=TRUE) # conditional=True, adjusts for correlations between predictors
```

```
#varimpAUC(cf1) # more robust towards class imbalance.
```

```
library(Boruta)
```

```
# Decide if a variable is important or not using Boruta
```

```
boruta_output <- Boruta( diagnosis~ ., data=na.omit(train_data), doTrace=2) # perform Boruta search
```

```
boruta_signif <- names(boruta_output$finalDecision[boruta_output$finalDecision %in%  
c("Confirmed", "Tentative")])
```

```
boruta_signif
```

```
#Model2: Naive Bayes
```

```
#Building and testing the model
```

```
model_nb <- train(diagnosis~.,  
  train_data,  
  method="nb",  
  metric="ROC",  
  preProcess=c('center', 'scale'),  
  trace=FALSE,  
  trControl=fitControl)
```

```
## predicting for test data
```

```
pred_nb <- predict(model_nb, test_data)
```

```
cm_nb <- confusionMatrix(pred_nb, test_data$diagnosis, positive = "M")
```

```
cm_nb
```

```
#Accuracy of the model is 93.9%
```

```
#Model3: glm
```

```
#Building and testing the model
```

```
model_glm <- train(diagnosis~.,  
  train_data,  
  method="glm",  
  metric="ROC",
```

```

preProcess=c('center', 'scale'),
trace=FALSE,
trControl=fitControl)

## predicting for test data
pred_glm <- predict(model_glm, test_data)
cm_glm <- confusionMatrix(pred_glm, test_data$diagnosis, positive = "M")
cm_glm
#Accuracy of the model is 98.3%
#algorithm for decision tree
library(C50)
data$diagnosis<-as.factor(data$diagnosis)
tree <- C5.0( diagnosis~., data = data)
summary(tree)
plot.new()
plot(tree)
results <- C5.0(diagnosis ~., data = data, rules = TRUE)
summary(results)
data<-as.data.frame(data)
library(rpart)
tree<-rpart(diagnosis~.,data =train_data,method="class")
plot(tree)
text(tree, pretty=0)
library(rattle)
library(rpart.plot)
library(RColorBrewer)
plot.new()
fancyRpartPlot(tree)
plot.new()
printcp(tree)
plotcp(tree)

```

```

ptree<- prune(tree, cp= tree$cptable[which.min(tree$cptable[, "xerror"]), "CP"])
plot.new()
fancyRpartPlot(ptree, uniform=TRUE,main="Pruned Classification Tree")
library(rpart)
fit1 <- rpart(diagnosis~.,data=train_data)
fit1
summary(fit1)
#Kernlab Classification
require(kernlab)
installed.packages("kernlab")
library(kernlab)
data_classifier<-ksvm(diagnosis ~., data =train_data , kernel='vanilladot')
data_classifier
data_predictions<-predict(data_classifier,test_data)
head(data_predictions)
table(data_predictions, test_data$diagnosis)
agreement<-data_predictions == test_data$diagnosis
table(agreement)
prop.table(table(agreement))
agreement
set.seed(12345)
data_classifier_rbf<-ksvm(diagnosis ~., data = train_data, kernel='rbfdot')
data_predictions_rbf<-predict(data_classifier_rbf,test_data)
agreement_rbf<-data_predictions_rbf == test_data$diagnosis
table(agreement_rbf)
prop.table(table(agreement_rbf))

# logistic regression model:
fit <- glm(diagnosis~.,data = train_data,family = binomial(link='logit'))
summary(fit)

```

```

library(MASS)

step_fit <- stepAIC(fit,method='backward')

summary(step_fit)

confint(step_fit)

#ANOVA on base model

anova(fit,test = 'Chisq')

#ANOVA from reduced model after applying the Step AIC

anova(step_fit,test = 'Chisq')


#plot the fitted model

plot.new()

plot(fit$fitted.values)

pred_link <- predict(fit,newdata = test_data,type = 'link')

#check for multicollinearity

library(car)

vif(fit)

vif(step_fit)

pred <- predict(fit,newdata =test_data ,type ='response')

#check the AUC curve

library(pROC)

g <- roc(diagnosis ~ pred, data = test_data)

g

plot.new()

plot(g)

library(caret)

#with default prob cut 0.50

test_data$pred_diagnosis <- ifelse(pred<0.5,'yes','no')


table(test_data$pred_diagnosis,test_data$diagnosis)


#training split of diagnosis classes

```



```

round(table(train_data$diagnosis)/nrow(train_data),2)*100

# test split of diagnosis
round(table(test_data$diagnosis)/nrow(test_data),2)*100

#predicted split of diagnosis
round(table(test_data$pred_diagnosis)/nrow(test_data),2)*100

#create confusion matrix
#confusionMatrix(test_data$diagnosis,test_data$pred_diagnosis)

#how do we create a cross validation scheme
control <- trainControl(method = 'repeatedcv',
                        number = 10,
                        repeats = 3)

seed <- 7
metric <- 'Accuracy'
set.seed(seed)

fit_default <- train(diagnosis~.,
                    data = train_data,
                    method = 'glm',
                    metric = metric ,
                    trControl = control)

print(fit_default)

library(caret)
varImp(step_fit)
varImp(fit_default)
library(woe)

library(riv)
train_data<-as.data.frame(train_data)
iv_df <- iv.mult(train_data, y="diagnosis", summary=TRUE, verbose=TRUE)
iv_df

iv <- iv.mult(train_data, y="diagnosis", summary=FALSE, verbose=TRUE)

# Plot information value summary

```

```
iv.plot.summary(iv_df)
```

```
#4. MARS (earth package)
```

```
#The earth package implements variable importance based on Generalized cross validation (GCV),
```

```
#number of subset models the variable occurs (nsubsets) and residual sum of squares (RSS).
```

```
library(earth)
```

```
marsModel<-earth(diagnosis~., data=data) # build model
```

```
ev <- evimp (marsModel) # estimate variable importance
```

```
ev
```

```
plot.new()
```

```
plot (ev)
```